LeanIMT: An optimized Incremental Merkle Tree

Privacy & Scaling Explorations

August 20, 2024

Abstract

This technical document presents the LeanIMT (Lean Incremental Merkle Tree), a data structure used to represent a group of elements efficiently. The LeanIMT is designed to optimize performance and reduce gas costs, making it suitable for zero-knowledge [11] protocols and applications.

Created: August 20, 2024

Updated: August 20, 2024

Contents

1 Introduction

This technical document aims to present and explain a new data structure called LeanIMT (Lean Incremental Merkle Tree). It covers the definition of Merkle Tree (MT), Incremental Merkle Tree (IMT) and Lean Incremental Merkle Tree (LeanIMT). It also explains the motivation behind creating this data structure and provides detailed algorithm explanations with pseudocode, correctness proofs and time complexity analyses. It also includes a section on benchmarks to illustrate performance improvements.

1.1 Motivation

The main motivation for the creation of this data structure was the development of the new version of the Semaphore protocol [5] [9] (version 4). Semaphore version 3 uses an IMT [2] [1], which is, however, rather inefficient and expensive when inserting the first leaves and when the number of leaves exceeds the maximum size supported by the tree.

2 Merkle Tree

A Merkle Tree (MT) is a tree (usually a binary tree) in which every leaf is a hash and every node that is not a leaf is the hash of its child nodes. [7] [4] [10]

2.1 Binary Tree

A Binary Tree is a tree data structure in which each node has at most two children, referred to as the left child and the right child. [6]

2.2 Incremental Merkle Tree

An Incremental Merkle Tree (IMT) is a Merkle Tree (MT) designed to be updated efficiently. [3]
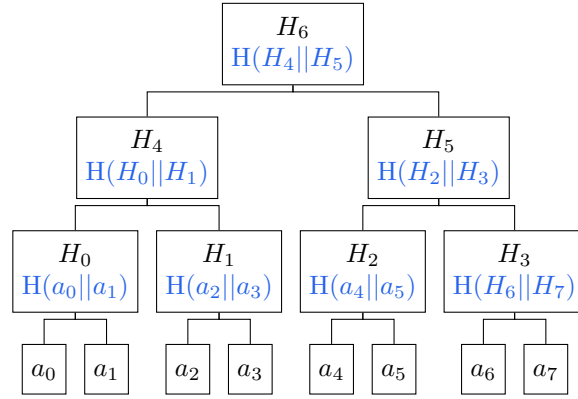
Example of IMT

3 LeanIMT

3.1 Definition

The **LeanIMT** (Lean Incremental Merkle Tree) is a Binary IMT.

The LeanIMT has two properties:

1. Every node with two children is the hash of its left and right nodes.

2. Every node with one child has the same value as its child node.

The tree is always built from the leaves to the root.

The tree will always be balanced by construction.

The tree depth is dynamic and can increase with the insertion of new leaves.

Example of a LeanIMT

T - Tree

V - Vertices (Nodes)

E - Edges (Lines connecting Nodes)

$T = (V, E)$

$V = \{a_0, a_1, a_2, H_0, H_1\}$

$E = \{(a_0, H_0), (a_1, H_0), (a_2, a_2), (H_0, H_1), (a_2, H_1)\}$

Example of LeanIMT

Note: To calculate a parent hash with two children, always start with the left child followed by the right. The order is never reversed.

3.2 Insertion

Function to insert a new leaf into a LeanIMT.

One of these cases will always be seen in each level when inserting a node:

1. The new node is the left child.

2. The new node is the right child.

Case 1: The new node is a left child

It will not be hashed, its value will be sent to the next level.

Adding a_4.

$T = (V, E)$

$V = \{a_0, a_1, a_2, a_3, H_0, H_1, H_2\}$

$E = \{(a_0, H_0), (a_1, H_0), (a_2, H_1), (a_3, H_1), (H_0, H_2), (H_1, H_2)\}$

Before inserting a_4   *After inserting a_4*

Case 2: The new node is a right child

The parent node will be the hash of the node's sibling with itself.

If we add a_5.

Before inserting a_5   *After inserting a_5*

3.2.1 Pseudocode

Algorithm 1 LeanIMT Insert algorithm

1: **procedure** INSERT(*leaf*)
2:    **if** $depth < newDepth$ **then**   ▷ $newDepth$ is the new depth of the tree after inserting the new node
3:   add a new empty array to nodes   ▷ Add a new tree level
4:    **end if**
5:    $node \leftarrow leaf$
6:    $index \leftarrow size$ ▷ The index of the new leaf equals the number of leaves in the tree.
7:    **for** level from 0 to depth - 1 **do**
8:       nodes[level][index] $\leftarrow$ node
9:       **if** index **is odd then**   ▷ It's a right node
10:          sibling $\leftarrow$ nodes[level][index - 1]
11:          node $\leftarrow$ **hash**(sibling, node)
12:       **end if**
13:       index $\leftarrow \lfloor index/2 \rfloor$ ▷ Divides the index by 2 and discards the remainder.
14:    **end for**
15:    nodes[depth] $\leftarrow$ [node]   ▷ Store the new root at the top level
16: **end procedure**

3.2.2 Correctness

Prove that after inserting a new leaf into the LeanIMT, the tree keeps all the properties.

The Mathematical Induction method will be used to prove the correctness of this algorithm.

n: Number of leaves in the tree.

Base Case

$n = 0$

If the tree is empty, the Insert function returns a new node with the value of the leaf. This satisfies the LeanIMT properties because the node does not have children. (Trivial case)

Inductive Hypothesis

Assume that with n leaves, the tree is a correct LeanIMT.

Inductive Step

$n \Rightarrow n + 1$

Assume that with n leaves, the tree is a correct LeanIMT and prove that with $n + 1$ leaves, the tree is still a correct LeanIMT.

The new node is always added at the end of the list of leaves.

To insert a new node, there are two possible cases:

1. When the new node is a left child.

2. When the new node is a right child.

Case 1: The new node is a left child

If the new node is a left child, it means that there were an even number of nodes at that level.

Then, since it's a left child, the parent has only one child and the parent has the same value as the child which is the new node. Then all the ancestors will be constructed following the two properties of the LeanIMT.

Since the tree was a correct LeanIMT with n nodes and inserting a new node that is a left child follows the properties of the LeanIMT $\Rightarrow$ the entire tree is still a correct LeanIMT.

Case 2: The new node is a right child

If the new node is a right child, it means that there were an odd number of nodes at that level.

Since it is a right child, the value of the parent will be the hash of the left child with the right child which is the new node. Then all the ancestors will be constructed following the two properties of the LeanIMT.

Since the tree was a correct LeanIMT with n nodes and inserting a new node that is a right child follows the properties of the LeanIMT $\Rightarrow$ the entire tree is still a correct LeanIMT.

Then for the two cases 1 and 2 the Insert algorithm follows the LeanIMT properties $\Rightarrow$ the Insert algorithm is correct.

Conclusion

Then by mathematical induction, the *insert* function works correctly.

3.2.3   Time complexity

n: Number of leaves in the tree.

d: Tree depth.

Every time a new node is added, it is necessary to update or add the ancestors up to the root of the tree.

Number of operations when adding a leaf: $\lceil d \rceil$

$\lceil d \rceil \leq d + 1$

$\qquad \leq O(\log n) + 1$

$\qquad \Rightarrow \boxed{O(\log n)}$

$d = \lceil \log(n) \rceil \leq \log(n) + 1$

$\qquad\qquad \Rightarrow O(\log n)$

The time complexity of the *Insert* function is $\boldsymbol{O(\log n)}$.

3.3   Batch Insertion

Performing the insertion in bulk rather than individually using a loop can lead to significant performance improvements because the number of hashing operations is reduced.

When inserting n members, all levels will be updated n times if the batch insertion function is not being used.

The core idea behind the batch insertion algorithm is to update each level only once even if there are many members to be inserted.

The algorithm will go through the nodes that are necessary to update the next level of the tree. The other nodes in the tree will not change.

Insert a_4 and a_5.

Before inserting a_4 and a_5

After inserting a_4 and a_5

3.3.1 Pseudocode

Algorithm 2 LeanIMT InsertMany algorithm

1: **procedure** INSERTMANY(*leaves*: List of nodes)
2:     $startIndex \leftarrow \lfloor size/2 \rfloor$ ▷ Divides the size of the tree by 2 and discards the remainder. startIndex is the index to start updating values at the next level.
3:     Add *leaves* to the tree leaves
4:     **for** level from 0 to depth - 1 **do**
5:         numberOfNodes $\leftarrow \lceil nodes[level].length/2 \rceil$ ▷ Calculate the number of nodes of the next level. numberOfNodes will be the smallest integer which is greater than or equal to the result of dividing the number of nodes of the level by 2.
6:         **for** index from *startIndex* to numberOfNodes - 1 **do**
7:             rightNode ← nodes[level][index * 2 + 1]   ▷ Get the right node if exists.
8:             leftNode ← nodes[level][index * 2]   ▷ Get the left node if exists.
9:             **if** *rightNode* exists **then**
10:                parentNode $\leftarrow hash(leftNode, rightNode)$
11:             **else**
12:                parentNode $\leftarrow leftNode$
13:   **end if**
14:             nodes[level + 1][index] ← parentNode ▷ Add the parent node to the tree.
15:         **end for**
16:         startIndex $\leftarrow \lfloor startIndex/2 \rfloor$ ▷ Divide startIndex by 2 and discards the remainder.
17:     **end for**
18: **end procedure**

3.3.2 Correctness

The idea is to prove that having a correct LeanIMT with n leaves, after using the *InsertMany* function to insert m new leaves, the resulting tree is also a correct LeanIMT.

Taking these LeanIMT properties as invariants:

1. Every node with two children is the hash of its left and right nodes.

2. Every node with one child has the same value as its child node.

The *InsertMany* algorithm will add all the new leaves and then calculate the

parent hashes. If a parent node has two children its value will be the hash of its left and right nodes, if it has one child, it will have the same value as its child.

Every time a new node is updated or added, all the invariants are maintained because the invariants are used to update/add nodes.

Then all the invariants are maintained before, during and after inserting m new leaves into a LeanIMT.

3.3.3 Time complexity

n: Number of leaves in the tree.

d: Tree depth.

m: Number of leaves to insert.

Number of operations when inserting elements in batch:

$\lceil m \rceil + \lceil \frac{m}{2} \rceil + \lceil \frac{m}{4} \rceil + \dots + \lceil \frac{m}{2^d} \rceil$

That is the same as $\sum_{k=0}^{d} \lceil \frac{m}{2^k} \rceil$

$\lceil m \rceil \leq m + 1$ then $\lceil \frac{m}{2^k} \rceil \leq \frac{m}{2^k} + 1$

$$\sum_{k=0}^{d} \lceil \frac{m}{2^k} \rceil \leq \sum_{k=0}^{d} (\frac{m}{2^k} + 1)$$

$$\leq \sum_{k=0}^{d} \frac{m}{2^k} + \sum_{k=0}^{d} 1$$

$$\leq 2m + O(\log(n + m))$$

$$\leq O(m) + O(\log(n + m))$$

$$\Rightarrow \boxed{O(m)}$$

$\sum_{k=0}^{d} \frac{m}{2^k} = m \sum_{k=0}^{d} \frac{1}{2^k} \approx m * 2 \Rightarrow 2m$

$\sum_{k=0}^{d} \frac{1}{2^k}$ (Geometric series)

$|r| < 1; r = \frac{1}{2}$

$$\sum_{k=0}^{\infty} a * r^k = \frac{a}{1-r} = \frac{1}{1-\frac{1}{2}}$$

$$= \frac{1}{\frac{1}{2}}$$

$$= 2$$

$$\sum_{k=0}^{d} 1 = d + 1$$

$$= O(\log(n + m)) + 1$$

$$\Rightarrow O(\log(n + m))$$

$$d = \lceil \log(n + m) \rceil \leq \log(n + m) + 1$$

$$\Rightarrow O(\log(n + m))$$

Then the time complexity of the *InsertMany* function is $\boldsymbol{O(m)}$.

Loop Insertion vs Batch Insertion

The time complexity of the Insertion function using a loop is $O(\log(n + m))$.

Going to the root to update or add nodes requires $\log(n + m)$ number of operations.

If going to the root to update or add nodes m times (one time per leaf to add) then the result will be:

$$m * \log(n + m) \Rightarrow \boxed{O(m \log(n + m))}$$

$\Rightarrow$ Time complexity Loop Insertion is superlinear: $\boldsymbol{O(m \log(n + m))}$

$\Rightarrow$ Time complexity Batch Insertion (*InsertMany* function) is linear: $\boldsymbol{O(m)}$

$\Rightarrow$ In terms of time complexity, it is more efficient to use the *InsertMany* function than the *Insert* function in a loop.

3.4   Update

Function to update the value of a leaf of a LeanIMT.

There are two cases:

1. When the node does not have a sibling.
2. When the node has a sibling.

Case 1: The node does not have a sibling

If the node that will be updated does not have a sibling, then the parent node will have the same value as the node.

Update a_4 to a_5

Before updating a_4 to a_5   *After updating a_4 to a_5*

Case 2: The node has a sibling

If the node has a sibling, the parent value will be the hash of the new value of the node with its sibling.

Update a_2 to a_5

Before updating a_2 to a_5   *After updating a_2 to a_5*

16

3.4.1 Pseudocode

Algorithm 3 LeanIMT Update algorithm

1: **procedure** UPDATE(*index*, *newLeaf*)
2:     $node \leftarrow newLeaf$
3:     **for** level from 0 to depth - 1 **do**
4:         nodes[level][index] ← node
5:         **if** index **is odd then**   ▷ It's a right node
6:             sibling ← nodes[level][index - 1]
7:             node ← **hash**(sibling, node)
8:         **else**   ▷ It's a left node
9:             sibling ← nodes[level][index + 1]
10:             **if** sibling exists **then**   ▷ It's a left node with a right sibling
11:                 node ← hash(node, sibling)
12:             **end if**
13:         **end if**
14:         index ← $\lfloor index/2 \rfloor$ ▷ Divides the index by 2 and discards the remainder.
15:     **end for**
16:     nodes[depth] ← [node]   ▷ Store the new root at the top level
17: **end procedure**

3.4.2 Correctness

Prove that after updating a node in the LeanIMT, the tree keeps all the properties.

The Mathematical Induction method will be used to prove the correctness of this algorithm.

The induction will be done using the depth of the tree. The correctness of the update function will be proven for all tree depths.

d: Depth of the tree.

Base Case

$d = 0$

If the depth is 0, it can have 0 or 1 node. If it has 0 leaves it is not necessary to use this function because there are no leaves. If it has 1 node, it is the root of the tree and if it is updated, the root is also updated because they are the same node.

Inductive Hypothesis

Assume that for a LeanIMT with depth d, the *update* function works correctly because after updating a leaf everything is updated correctly up to the root.

Inductive Step

$d \Rightarrow d + 1$

If for a LeanIMT with depth d, the *update* function works correctly because after updating a leaf everything is updated correctly up to the root, then prove that for depth $d + 1$ the update function also works correctly.

There are two cases when updating a leaf:

1. When the node does not have a sibling.

2. When the node has a sibling.

Case 1: The node does not have a sibling

If the node that will be updated does not have a sibling, then the parent node will have the same value as the node.

**Case 2: The node has a sibling**

If the node has a sibling, the parent value will be the hash of the node with its sibling.

After those cases, the tree that should be updated has depth d. Using the inductive hypothesis, the subtree of depth d will correctly update all the ancestors up to the root.

Conclusion

Then by mathematical induction, the update function works correctly for any depth of the tree.

3.4.3   Time complexity

n: Number of leaves in the tree.

d: Tree depth.

Every time a leaf is updated, it is necessary to update all the ancestors up to the root of the tree.

Number of operations when updating a leaf: $\lceil d \rceil$

This proof is the same as the proof of the time complexity of the Insert function.

The time complexity of the *Update* function is $O(\log n)$.

3.5   Remove

Function to remove a leaf from a LeanIMT.

The *remove* function is the same as the *update* function but the value used to update is 0.

You can use a value other than 0, the idea is to use a value that is not a possible value for a correct member in the list.

Remove a_2

Before removing a_2   *After removing a_2*

3.5.1   Pseudocode

Algorithm 4 LeanIMT Remove algorithm

1: **procedure** REMOVE(*index*)
2:     *update*(*index*, 0)
3: **end procedure**

3.5.2   Correctness

The proof of the correctness of this algorithm is the same as the Update function.

19

3.5.3 Time complexity

The proof of the time complexity of this algorithm is the same as the Update function.

3.6 Generate Merkle Proof

Function to generate a Merkle Proof of a leaf in a LeanIMT.

There are two cases:

1. When the node does not have a sibling.

2. When the node has a sibling.

d: Depth of the tree

Case 1: The node does not have a sibling

When generating the proof for this case, nothing is added to the proof at that level.

This case only happens when the node is the last node in the level and is also a left node.

If we want to generate a proof for the node a_4.

LeanIMT to generate a proof for a_4   *Nodes used to generate a proof for a_4*

path: [1]

Merkle Proof: {

root: H_3

leaf: a_4

index: 1

siblings: $[H_2]$

}

Case 2: The node has a sibling

When generating the proof for this case, if the node is a right node, 1 will be added to the path and the left sibling will be added to siblings and if the node is a left node, 0 will be added to the path and the right sibling will be added to siblings.

If we want to generate a proof for the node a_3.
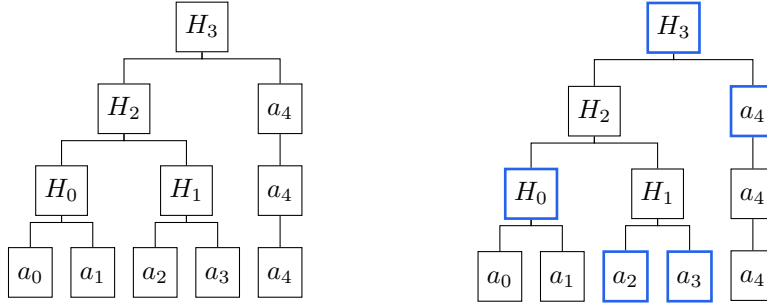
LeanIMT to generate a proof for a_3   *Nodes used to generate a proof for a_3*

path: $[1, 1, 0]$

Merkle Proof: {

root: H_3

leaf: a_3

index: 3

siblings: $[a_2, H_0, a_4]$

}

3.6.1 Pseudocode

Algorithm 5 LeanIMT generateProof algorithm

1: **procedure** GENERATEPROOF(*index*)
2:     siblings ← empty list ▷ List to store the nodes necessary to rebuild the root.
3:     path ← empty list   ▷ List of 0s or 1s to help rebuild the root. 0 if the current node is a left node with a sibling and 1 if the current node is a right node.
4:     **for** level from 0 to depth - 1 **do**
5:         isRightNode ← index is odd
6:         **if** isRightNode is true **then**   ▷ It's a right node
7:             siblingIndex ← index - 1
8:         **else**   ▷ It's a left node
9:             siblingIndex ← index + 1
10:         **end if**
11:         sibling ← nodes[level][siblingIndex]
12:         **if** *sibling* exists **then**
13:             add isRightNode to path
14:             add sibling to siblings
15:         **end if**
16:         index ← $\lfloor index/2 \rfloor$ ▷ Divides the index by 2 and discards the remainder.
17:     **end for**
18:     leaf ← leaves[index]
19:     index ← reverse path and use the list as a binary number and get the decimal representation
20:     siblings ← leaves[index]
21:     proof ← $\{root, leaf, index, siblings\}$
22:     **return** proof
23: **end procedure**

Note:

The *index* parameter of the function and the *index* returned in the proof are different variables that share the same name but may not always have same value. The *index* in the parameter is the index of the element that will be used to generate the proof and the *index* in the proof is the decimal representation of the reverse of the path.

If the number of elements in the path is equal to the depth of the tree, it

means that all levels (except the root level) were included in the proof because they are necessary to calculate the root. In this case, the *index* in the proof will have the same value as the *index* in the function parameter.

However, if the path contains fewer elements than the depth of the tree, it means that some levels (except the root level) were omitted from the proof because they were unnecessary for calculating the root value. In this case, the two *index* values will differ.

3.6.2 Correctness

Prove that the *generateProof* function works correctly for all tree depths.

The Mathematical Induction method will be used to prove the correctness of this algorithm.

The induction will be done using the depth of the tree. The correctness of the *generateProof* function will be proven for all tree depths.

d: Depth of the tree.

Base Case

$d = 0$

If the depth is 0, it can have 0 or 1 node. If it has 0 leaves it is not necessary to use this function because there are no leaves. If it has 1 leaf, the merkle proof will have the root and the leaf with the same values because they are the same nodes. The siblings will be empty and the path (index) too.

Inductive Hypothesis

Assume that for a LeanIMT with depth d, the *generateProof* function works correctly because it returns the correct values necessary to rebuild the root.

Inductive Step

$d \Rightarrow d + 1$

If for a LeanIMT with depth d, the *generateProof* function works correctly, prove that for a tree with depth $d + 1$, the *generatProof* function also works correctly.

There are two cases when generating a proof:

1. When the node does not have a sibling.

2. When the node has a sibling.

Case 1: The node does not have a sibling

If the node does not have a sibling, nothing is added to the proof.

Case 2: The node has a sibling

If the node has a sibling, both the path and siblings arrays are updated accordingly. If the node is a right node, 1 will be added to the path and the left sibling will be added to siblings. If the node is a left node, 0 will be added to the path and the right sibling will be added to siblings.

After those cases, the tree that should generate the proof has depth d. Using the inductive hypothesis, the subtree of depth d will correctly return the values in the proof.

Conclusions

By mathematical induction, the *generateProof* function works correctly for any depth of the tree.

3.6.3 Time complexity

n: Number of leaves in the tree.

d: Tree depth.

To generate a Merkle Proof it is necessary to visit all the ancestors of the leaf up to the root of the tree.

Number of operations to generate a Merkle Proof: $\lceil d \rceil$

This proof is the same as the proof of the time complexity of the Insert function.

The time complexity of the *generateProof* function is $\boldsymbol{O(\log n)}$.

3.7 Verify Merkle Proof

Function to verify a Merkle Proof of a leaf in a LeanIMT.

The *verifyProof* function will verify if a leaf is part of a tree having a Merkle Proof.

The algorithm will go through the sibling nodes using the path and calculate the parent in the next level of the tree. Then it will check if the calculated root matches the one that is part of the proof. If the calculated root matches the one that is part of the root, the algorithm will return true, otherwise it will return false.

Diagrams when verifying a Merkle Proof

Verifying a proof for node a_4

Note: This proof was generated using case 1 of the *generateProof* function.

path: [1]

Merkle Proof: {

root: H_3

leaf: a_4

index: 1

siblings: $[H_2]$

}

LeanIMT to verify a proof for a_4   *Nodes rebuilt to verify a proof for a_4*

Verifying a proof for node a_3

Note: This proof was generated using case 2 of the *generateProof* function.

25

path: [1, 1, 0]

Merkle Proof: {

root: H_3

leaf: a_3

index: 3

siblings: $[a_2, H_0, a_4]$

}

LeanIMT to verify a proof for a_3   *Nodes rebuilt to verify a proof for a_3*

3.7.1 Pseudocode

Algorithm 6 LeanIMT verifyProof algorithm

1: **procedure** VERIFYPROOF($index$)
2:     { root, leaf, siblings, index } ← proof   ▷ Deconstruct the proof
3:     node ← leaf
4:     **for** i from 0 to siblings.length - 1 **do**
5:         isOdd ← divide index by 2 i times and check if the result is odd
6:         **if** isOdd is true **then**   ▷ *node* is a right child
7:             node ← $hash(siblings[i], node)$
8:         **else**   ▷ It's a left node
9:             node ← $hash(node, siblings[i])$
10:         **end if**
11:     **end for**
12:     **if** *root* is equal *node* **then**
13:         **return** true
14:     **else**
15:         **return** false
16:     **end if**
17: **end procedure**

3.7.2 Correctness

Prove that the *verifyProof* function works correctly for all tree depths.

The Mathematical Induction method will be used to prove the correctness of this algorithm.

The induction will be done using the depth of the tree. The correctness of the *verifyProof* function will be proven for all tree depths.

d: Depth of the tree.

Base Case

$d = 0$

If the depth is 0, it can have 0 or 1 node. If it has 0 leaves it is not necessary to use this function because there are no leaves. If it has 1 leaf, the merkle root will be equal to the node so it will verify the proof successfully.

Inductive Hypothesis

Assume that for a LeanIMT with depth d, the *verifyProof* function works correctly because it knows how to calculate all the hashes up to the root and

then check if the calculated hash is equal to the root in the proof.

Inductive Step

$d \Rightarrow d + 1$

If for a LeanIMT with depth d, the $verifyProof$ function works correctly, prove that for a tree with depth $d + 1$, the $verifyProof$ function also works correctly.

If the current node (which is the leaf at the beginning of the algorithm) is a right node, the parent value will be the hash of the left sibling with the current node. If the current node is a left node, the parent value will be the hash of the current node with the right sibling. Then, the current node will be the parent node.

After running that once, the tree that should verify the proof has depth d. Using the inductive hypothesis, the subtree of depth d will successfully verify the values of the proof.

Conclusion

Then the $verifyProof$ function works correctly for any depth of the tree.

3.7.3 Time complexity

n: Number of leaves in the tree.

d: Tree depth.

To verify a Merkle Proof it is necessary to visit (rebuild) all the ancestors of the leaf up to the root of the tree.

Number of operations to verify a Merkle Proof: $\lceil d \rceil$

This proof is the same as the proof of the time complexity of the Insert function.

The time complexity of the $verifyProof$ function is $\boldsymbol{O(\log n)}$.

4 Implementations

The TypeScript/JavaScript and Solidity implementations follow the same idea and are compatible but are different.

The TypeScript/JavaScript implementation focuses on performance whereas the Solidity one focuses on saving gas costs.

The TypeScript/JavaScript and Solidity code of the LeanIMT was audited as part of the Semaphore v4 audit [8].

4.1   TypeScript/JavaScript

TypeScript/JavaScript LeanIMT code: https://github.com/privacy-scaling-explorations/zk-kit/tree/main/packages/lean-imt

4.2   Solidity

Solidity LeanIMT code:
https://github.com/privacy-scaling-explorations/zk-kit.solidity/tree/main/packages/lean-imt

5   Benchmarks

All the benchmarks were run in an environment with these properties:

System Specifications

Computer: MacBook Pro

Chip: Apple M2 Pro

Memory (RAM): 16 GB

Operating System: macOS Sonoma version 14.5

Software environment

Node.js version: 20.5.1

Browser: Google Chrome Version 127.0.6533.73 (Official Build) (arm64)

5.1   Running the benchmarks

TypeScript/JavaScript

GitHub repository to run Node.js and browser benchmarks:
https://github.com/vplasencia/imt-benchmarks.

Solidity

GitHub repository to run Solidity benchmarks:
https://github.com/privacy-scaling-explorations/zk-kit.solidity

5.2 TypeScript/JavaScript

Note: The following TypeScript/JavaScript benchmarks demonstrate how to perform the same operations for each data structure. This means that even if the remove function of the IMT is called *Delete* in the implementation, it will be referred to as *Remove* in the benchmarks. Additionally, the IMT does not have a function called $InsertMany$, the $IMT - InsertMany$ benchmarks use the *Insert* function in a loop to simulate inserting multiple members at the same time.

Note: Although the IMT does not have a batch insertion function, it does include an optimization for adding multiple members through its constructor.

Note: The IMT has a static depth. To run the benchmarks, the minimum depth necessary to perform the operation was used.

For example:

- If the IMT has 4 members and I want to add 1 new member the tree depth used will be 3.

- If the IMT has 5 members and I want to add 1 new member the tree depth used will be 3.

5.2.1 Node.js

Table 1: All Functions (100 iterations)

Function	ops/sec	Average Time (ms)	Relative to IMT
IMT - Insert	1287	0.77687	
LeanIMT - Insert	2358	0.42391	1.83 x faster
IMT - InsertMany	12	77.98467	
LeanIMT - InsertMany	144	6.94025	11.24 x faster
IMT - Update	1283	0.77933	
LeanIMT - Update	1223	0.81708	1.05 x slower
IMT - Remove	1306	0.76554	
LeanIMT - Remove	1301	0.76838	1.00 x slower
IMT - GenerateProof	300868	0.00332	
LeanIMT - GenerateProof	321586	0.00311	1.07 x faster
IMT - VerifyProof	1331	0.75121	
LeanIMT - VerifyProof	1336	0.74810	1.00 x faster

Figure 1: Functions IMT vs LeanIMT (100 iterations)

Figure 2: Functions IMT vs LeanIMT (100 iterations)

5.2.2 Browser

Table 2: All Functions (100 iterations)

Function	ops/sec	Average Time (ms)	Relative to IMT
IMT - Insert	1107	0.90300	
LeanIMT - Insert	2590	0.38600	2.34 x faster
IMT - InsertMany	14	68.53200	
LeanIMT - InsertMany	158	6.30200	10.87 x faster
IMT - Update	1455	0.68700	
LeanIMT - Update	1470	0.68000	1.01 x faster
IMT - Remove	1438	0.69500	
LeanIMT - Remove	1472	0.67900	1.02 x faster
IMT - GenerateProof	1000000	0.00100	
LeanIMT - GenerateProof	1000000	0.00100	1.00 x slower
IMT - VerifyProof	1472	0.67900	
LeanIMT - VerifyProof	1508	0.66300	1.02 x faster

Figure 3: Functions IMT vs LeanIMT (100 iterations)

Figure 4: Functions IMT vs LeanIMT (100 iterations)

5.2.3 LeanIMT: Node.js vs Browser

Figure 5: LeanIMT Node.js vs Browser (100 iterations)

5.2.4 Insert Function: IMT vs LeanIMT

Table 3: Insert Function (1000 iterations)

Function	ops/sec	Average Time (ms)	Relative to IMT
IMT	814	1.22803	
LeanIMT	1453	0.68790	1.79 x faster

Figure 6: Insert function IMT vs LeanIMT (1000 iterations)

Figure 7: Insert function IMT vs LeanIMT (1000 iterations)

5.2.5 LeanIMT: Insert Loop vs Batch Insertion

Comparing the LeanIMT *Insert* function in a loop with the LeanIMT *InsertMany* function.

Table 4: Insert Function (100 iterations)

Function	ops/sec	Average Time (ms)	Relative to Insert
Insert in Loop	47	20.97820	
InsertMany	136	7.31698	2.87 x faster

Figure 8: Batch Insertion LeanIMT (100 iterations)

Figure 9: Batch Insertion LeanIMT (100 iterations)

5.3 Solidity

```
│ Solidity and Network Configuration                                                               │
│·······························|·············|··············|··············|···················· │
│ Solidity: 0.8.23             · Optim: true · Runs: 200    · viaIR: false ·    Block: 30,000,000 gas │
│ Methods                                                                                          │
│·······························|·············|··············|··············|··············|········ │
│ Contracts / Methods          · Min         · Max          · Avg          · # calls     · usd (avg) │
│ BinaryIMTTest                ·                                                                    │
│······························|·············|··············|··············|··············|········ │
│     init                     ·     105,471 ·      374,307 ·      357,505 ·          16 ·        – │
│······························|·············|··············|··············|··············|········ │
│     initWithDefaultZeroes    ·      91,272 ·       91,870 ·       91,471 ·           3 ·        – │
│······························|·············|··············|··············|··············|········ │
│     insert                   ·      98,112 ·    2,501,619 ·      560,351 ·          31 ·        – │
│······························|·············|··············|··············|··············|········ │
│     remove                   ·     471,034 ·      473,216 ·      472,710 ·           7 ·        – │
│······························|·············|··············|··············|··············|········ │
│     update                   ·           – ·            – ·      474,000 ·           1 ·        – │
│ Deployments                  ·                                            · % of limit           │
│······························|·············|··············|··············|··············|········ │
│ BinaryIMT                    ·   1,237,933 ·    1,238,005 ·    1,237,998 ·       4.1 % ·        – │
│······························|·············|··············|··············|··············|········ │
│ BinaryIMTTest                ·     378,277 ·      378,337 ·      378,329 ·       1.3 % ·        – │
│······························|·············|··············|··············|··············|········ │
│ PoseidonT3                   ·           – ·            – ·    3,693,362 ·      12.3 % ·        – │
│ Key                                                                                              │
│··················································································· │
│ ○  Execution gas for this method does not include intrinsic gas overhead                         │
│··················································································· │
│ △  Cost was non-zero but below the precision setting for the currency display (see options)      │
│··················································································· │
│ Toolchain:  hardhat                                                                              │
```

Figure 10: IMT Gas Report

```
│ Solidity and Network Configuration                                                      │
│ ···········································································              │
│ Solidity: 0.8.23  ·  Optim: true  ·  Runs: 200  ·  viaIR: false  ·  Block: 30,000,000 gas │
│ ···········································································              │
│ Methods                                                                                 │
│ ···········································································              │
│ Contracts / Methods  ·  Min        ·  Max        ·  Avg        ·  # calls    ·  usd (avg) │
│ ···········································································              │
│ LeanIMTTest                                                                             │
│ ···········································································              │
│     insert         ·      93,938  ·     163,708  ·    119,051  ·         47  ·         — │
│ ···········································································              │
│     insertMany     ·      95,891  ·     715,164  ·    322,619  ·          7  ·         — │
│ ···········································································              │
│     remove         ·     104,558  ·     296,279  ·    233,235  ·         13  ·         — │
│ ···········································································              │
│     update         ·      58,909  ·     252,738  ·    197,830  ·          8  ·         — │
│ ···········································································              │
│ Deployments                                                                 % of limit  │
│ ···········································································              │
│ LeanIMT            ·   1,018,010  ·   1,018,082  ·  1,018,077  ·      3.4 %  ·         — │
│ ···········································································              │
│ LeanIMTTest        ·     455,827  ·     455,911  ·    455,908  ·      1.5 %  ·         — │
│ ···········································································              │
│ PoseidonT3         ·          —   ·          —   ·  3,693,362  ·     12.3 %  ·         — │
│ ···········································································              │
│ Key                                                                                     │
│ ···········································································              │
│ ○ Execution gas for this method does not include intrinsic gas overhead                 │
│ ···········································································              │
│ △ Cost was non-zero but below the precision setting for the currency display (see options) │
│ ···········································································              │
│ Toolchain:  hardhat                                                                     │
```

Figure 11: LeanIMT Gas Report

Figure 12: Gas cost of the execution of the Functions IMT vs LeanIMT

6 Conclusions

This technical document explains the LeanIMT algorithms, proves their correctness and analyzes their time complexity. The benchmarks show the improvements of LeanIMT, which is the data structure used in Semaphore v4, over the IMT used in Semaphore v3.

6.1 Future Work

As future work, a function to update many members at once (similar to the *insertMany* function to insert many members at once) will be developed. Additionally, a Rust implementation of the data structure will be created to benchmark the performance in Node.js and browser environments.

References

[1] *IMT Solidity implementation*. URL: https://github.com/privacy-scaling-explorations/zk-kit.solidity/tree/main/packages/imt.

[2] *IMT TypeScript implementation*. URL: https://github.com/privacy-scaling-explorations/zk-kit/tree/main/packages/imt.

[3] *Incremental Merkle Tree*. URL: https://zokyo-auditing-tutorials.gitbook.io/zokyo-tutorials/tutorial-16-zero-knowledge-zk/definitions-and-essentials/incremental-merkle-tree.

[4] G. Mohr J. Chapweske. "Tree Hash EXchange format (THEX)". In: (2003). URL: https://web.archive.org/web/20090803220648/http://open-content.net/specs/draft-jchapweske-thex-02.html.

[5] Barry Whitehat Kobi Gurkan Koh Wei Jie. "Semaphore: Zero-Knowledge Signaling on Ethereum". In: (2020). URL: https://semaphore.pse.dev/whitepaper-v1.pdf.

[6] NIST. "Binary Tree". In: (2017). URL: https://xlinux.nist.gov/dads/HTML/binarytree.html.

[7] NIST. "Merkle Tree". In: (2019). URL: https://xlinux.nist.gov/dads/HTML/MerkleTree.html.

[8] PSE. "Semaphore v4 Audit Report". In: (2024). URL: https://semaphore.pse.dev/Semaphore_4.0.0_Audit.pdf.

[9] *Semaphore Website*. URL: https://semaphore.pse.dev/.

[10] Alin Tomescu. "What is a Merkle Tree?" In: (2020). URL: https://decentralizedthoughts.github.io/2020-12-22-what-is-a-merkle-tree/.

[11]   *What are zero-knowledge proofs?* URL:
       https://ethereum.org/en/zero-knowledge-proofs/.